

same scoring, every system:

bare LLM 0–12% usable · hallucination ~1.0

Labwright 88–100% usable · hallucination 0.000

goal
(natural language)

**LLM agent
proposes raw inputs**

**deterministic
calculators (physics)**

**verifier
re-proves
every number**

**SOP + design JSON
verified**

hard gate:
submit_design

invalid design → agent retries

the model never writes a derived number

every derived number is recomputed from the model's own raw inputs

the alternatives an LLM reaches for — rejected:

type the
numbers

soft self-check
(LLM judges its own work)

REJECTED

a number you typed is not trusted

honest boundary: the gate verifies
arithmetic, not physiology — on blind goals
usable drops, hallucination stays 0.000